

1001 Sherbrooke St. W

Montréal, Canada

✉ li.jiang3@mail.mcgill.ca

🌐 louieworth.github.io

🔗 eBZfztYAAAAJ&hl

Li Jiang

Education

2023– **McGill University**, Montréal, Canada
Mila – Québec AI Institute, Montréal, Canada
Ph.D. in Operation Management
Advisors: Prof. Yichuan Ding

2020–2023 **Tsinghua University**, Beijing, China
M.S in Computer Science
Thesis: Offline Reinforcement Learning with In-Sample Paradigm
Advisors: Prof. Xianyuan Zhan and Prof. Wai Kin (Victor) Chan

2014–2018 **Southwest University of Sci. & Tech.**, Sichuan, China
B.E in Civil Engineering

Research Experiences

2022.8–2023.1 **CMU**, Pittsburgh, Visiting Student
Worked on safety in offline reinforcement learning with Prof. Ding Zhao.

2022.6–2022.8 **Stanford University**, Palo Alto, Student Research Fellow
Worked on AI safety and alignment.

2021.6–2022.6 **Tsinghua University**, Beijing, Research Intern
Worked on offline reinforcement learning and 5G MMIMO optimization with Prof. Xianyuan Zhan.

2021.1–2021.6 **Microsoft Research Asia**, Beijing, Research Intern
Worked on reinforcement learning for the molecular generation with Prof. Jiang Bian.

Selected Publications

See all publications on [Google Scholar](#) (* indicates Equal Contribution).

1. **Hummer: Towards Limited Competitive Preference Dataset**
L. Jiang*, Y. Wu*, J. Xiong, J. Ruan, Y. Ding, Q. Guo, Z. Wen, J. Zhou, X. Deng
[\[ArXiv\]](#)
Conference on Language Models (CoLM), 2024.
ICML Workshop on Models of Human Feedback for AI Alignment (Oral), 2024.
ICML Workshop on the Next Generation of AI Safety, 2024.
2. **Offline RL with No OOD Actions: Offline RL with Implicit Value Regularization** [\[Paper\]](#)
H. Xu, L. Jiang, J. Li, Z. Yang, Z. Wang, V.W.K Chan, X. Zhan
International Conference on Learning Representations (ICLR), 2023 (Notable Top 5% in all accepted papers).

3. **A Policy-Guided Imitation Approach for Offline Reinforcement Learning** [\[Paper\]](#)
L. Jiang*, H. Xu*, J. Li, X. Zhan
Neural Information Processing Systems (NeurIPS), 2022 (Oral < 2% in all accepted papers).
4. **Curriculum Goal-conditioned Imitation for Offline Reinforcement Learning** [\[Paper\]](#)
L. Jiang*, X. Feng*, X. Yu, H. Xu, X. Sun, X. Zhan, V.W.K Chan
IEEE Transactions on Games (ToG), 2022.

Selected Pre-prints

1. **Which heads matter for reasoning? rl-guided kv cache compression**
W. Du, **L. Jiang***, K. Tao, X. Liu, H. Wang. [\[ArXiv\]](#)
Latent & Implicit Thinking Workshop, ICLR 2026.
2. **Beyond Single-Shot: A Multi-Sample Distributional Framework for Consumer Preference Alignment**
L. Jiang*, X. Liu. Working paper.

Awards and Fellowships

- 2024 OpenAI's Researcher Access Program (USD \$1,000). [\[Detail\]](#)
- 2023 Outstanding reviewer at NeurIPS.
- 2023 Shenzhen Universiade International Scholarship Foundation (CAD \$11,000). [\[Detail\]](#)
- 2023 Excellent Graduate Student at Tsinghua University (2%).
- 2022 Stanford Summer Research Fellowship (CAD \$10,169). [\[Detail\]](#)
- 2022 Tsinghua University Second-class Scholarship (CAD \$2,000, 5%).
- 2018 Scholarship of China Scholarship Council (CAD \$3,000, among 30 recipients in China.)

Invited Talks and Presentations

- Feb. 2023 **Offline RL with Implicit Value Regularization**, *CMU Safe AI Lab*, Host: Prof. Ding Zhao.
- Nov. 2022 **A Policy-Guided Imitation Approach for Offline RL**, *The 4th TBSI Workshop on Learning Theory (TBSIWOLT'22)*, hosted by Prof. Yang Li.
- Aug. 2022 **A Policy-Guided Imitation Approach for Offline RL**, *RL China*.

Professional Service

- Teaching Assistant
 - 2023 Fall MGCR 372, Operation Management
 - 2024 Winter MGCR 372, Operation Management
- Conference and Journal Reviewing
 - Reviewer for JMLR, TMLR, ICLR (2023-2025), NeurIPS (2022-2024), ICML (2023)

(Latest updated on March 2, 2026)